

# Part 1. Classical Approaches To Recommendation

Olivier Koch

Criteo Research

*June 28, 2018*

# Overview

1. Neighborhood methods
2. Matrix Factorization
3. Latent factor models
4. Hybrid methods

# Problem statement

- Item recommendation with implicit or explicit feedback data
- Given interactions between users and items, find the  $k$  most relevant items for a given user.

## More formally...

Given:

- $U = (u_1, \dots, u_n)$ , a group of  $n$  users
- $I = (i_1, \dots, i_m)$ , a group of  $m$  items
- $R_{u,i}$ , a set of interactions between users and items

Find the  $k$  most relevant items for each user

# Overview of methods

- Neighborhood-based methods
- Matrix Factorization
- Other latent factor models
- Hybrid methods

# Lessons from the Netflix Prize Challenge

- Open competition to beat the Netflix baseline (CineMatch)
- Training set: 100,480,507 ratings (480,189 users, 17,770 movies)
- Qualifying set: 2,817,131 ratings (50% test set, 50% quiz set)
- Improve by 10% to win \$1,000,000

# Lessons from the Netflix Prize Challenge

- A trivial solution is almost as good as the baseline (average all ratings for a movie)
- Use ensemble methods
- In particular, combine neighborhood-based methods and MF
- Incorporate side information (e.g. time of rating)

# Neighborhood methods

- User-oriented: find similar users
- Item-oriented: find similar items



## Item-based recommendation

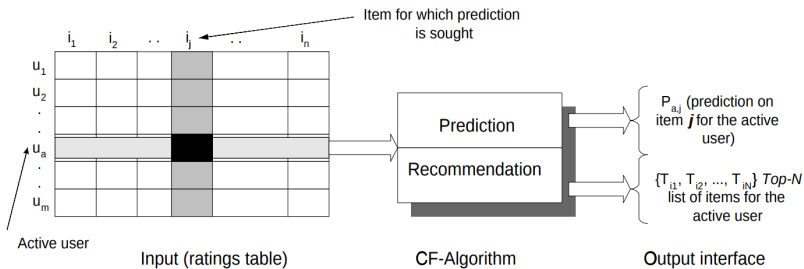
To predict the rating of  $u_i$  and  $i_j$ :

- Find the  $k$  most similar items to  $u_i$
- Keep items that  $i_j$  has rated
- Weight each rating by the similarity between these items and  $i_j$

$$P_{u,i} = \frac{\sum_{all\ similar\ items} s_{i,j} \times R_{u,j}}{\sum_{all\ similar\ items} s_{i,j}}$$

Item-based top-n recommendation algorithms, M.Deshpande and G. Karypis.  
Trans. Inf. Syst. 2004.

# Item-based recommendation



Item-based collaborative filtering recommendation algorithms, B.M. Sarwar, G. Karypis, J.A. Konstan, and J. Reidl. WWW'2001.

# Similarity definition

Cosine-based similarity

$$s_{i,j} = \cos(\vec{i}, \vec{j}) = \frac{\vec{i} \cdot \vec{j}}{\|\vec{i}\| \cdot \|\vec{j}\|}$$

Pearson similarity

$$s_{i,j} = \frac{\sum_{u \in U} (R_{u,i} - \bar{R}_i)(R_{u,j} - \bar{R}_j)}{\sqrt{\sum_{u \in U} (R_{u,i} - \bar{R}_i)^2} \sqrt{\sum_{u \in U} (R_{u,j} - \bar{R}_j)^2}}$$

Adjusted cosine similarity

$$s_{i,j} = \frac{\sum_{u \in U} (R_{u,i} - \bar{R}_u)(R_{u,j} - \bar{R}_u)}{\sqrt{\sum_{u \in U} (R_{u,i} - \bar{R}_u)^2} \sqrt{\sum_{u \in U} (R_{u,j} - \bar{R}_u)^2}}$$

Let's get hands-on...

Exercise 1: item-based collaborative filtering in Python

# Neighborhood vs Matrix Factorization

- Neighborhood models detect localized relationships
- MF captures the totality of weak signals contained in the matrix
- Both methods can be fused (e.g. SVD++)
- Both methods have trouble making predictions for users with few ratings

Factorization Meets the Neighborhood: a Multifaceted Collaborative Filtering Model, Y. Koren, KDD'08

# SVD vs Matrix Factorization

- SVD (Singular Value Decomposition) is closed form, requires a full-filled matrix (biased towards unseen observations)
- Matrix Factorization uses optimization and has been shown to work better (e.g. on known entries only)

# Singular Value Decomposition

$$\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^*$$

where  $\mathbf{U}$  and  $\mathbf{V}$  contain the left- and right-singular vectors of  $\mathbf{M}$ .

# Matrix factorization

- Find a low rank approximation to a sparse ratings matrix by minimizing a loss function on the known ratings
- Popular choice for applications in industry
- Simple
- Highly scalable to large datasets
- Captures global signals (but fails to capture strong local signals)



More formally...

$$R_{u,i} = U_{n \times k}^T * V_{k \times m}$$

where  $U_{n \times k}$ : user matrix and  $V_{k \times m}$  item matrix

# Steps for matrix factorization

- Build user-item matrix
- Define factorization model (cost function)
- Remove global mean
- Remove item bias and user bias
- Add regularization to prevent overfitting
- Minimize cost function (e.g. stochastic gradient descent, alternating least squares)

## More formally...

Find  $U$  and  $V$  such that:

$$\min_{U, V} \sum_{(u, i) \in k} (R_{u, i} - \mu - b_i - b_u - U_i^T V_u)^2 + \lambda (\|U_i\|^2 + \|V_u\|^2 + b_i^2 + b_u^2)$$

# Matrix Factorizations parameters

- cost function
- vector size
- learning rate
- regularization
- optimization method

Let's get hands-on...

Exercise 2: MF with tensorflow

# Matrix Factorization for Implicit Feedback

- Model the probability that a user will interact with a given item with a logistic function
- Weigh zero-entries
- Optimize log likelihood with Alternating gradient descent
- Outperforms regular MF on small vectors (Spotify dataset)

# Probabilistic Matrix Factorization

- Scales linearly with number of observations
- Can incorporate the fact that users who rate similar sets of movies have similar preferences
- Evaluation on the Netflix dataset

$$p(R \mid q, p, \sigma^2) = \prod_{i=1}^N \prod_{j=1}^M [\mathcal{N}(R_{ij} \mid U_i^T V_j, \sigma^2)]^{I_{ij}}$$


# Latent factor models

- Probabilistic Latent Semantic Indexing (PLSA)
- Latent Dirichlet Allocation (LDA)
- Restricted Boltzmann Machines (RBM)



## Contextual side information


- High-order SVD optimized on observed entries only
- Batch sub-space descent  $\rightarrow$  SGD for larger datasets



Multiverse Recommendation: N-dimensional Tensor Factorization for Context Aware Collaborative Filtering, Karatzoglou A. Amatriain X. Baltrunas L. Oliver N. RecSys 2010

# Temporal dynamics

- Including temporal dynamics in user and item behaviors
- Applicable to neighborhood-based and factorization methods



# Unified Boltzmann machines

- Model the joint distribution of a set of binary variables through their pairwise interactions
- Encode collaborative and content information as features
- Learn weights that reflect how well each feature predicts user actions

A Unified Approach to Building Hybrid Recommender Systems, Gunawardana A.  
Meek C., RecSys 2009

# Lessons from the Real World

- What problem are you trying to solve? Predicting DVD rentals is different from predicting instant streaming selection.
- Pick a single solution (one that fits your problem)
- There is often a deceptively simple baseline solution
- Improve on it incrementally
- Evaluate online early

# Exercise 1: item-based recommendation with Movie Lens

Fetch the code and data:

<https://github.com/oakfr/intro-to-reco>

## Exercise 1: item-based recommendation with Movie Lens

```
from sklearn.metrics import pairwise_distances
from scipy.spatial.distance import cosine, correlation

def compute_movies_similarities (method='cosine'):
    [...]
    movie_sim = # PUT YOUR CODE HERE
    return pd.DataFrame( movie_sim )
```

## Exercise 1: item-based recommendation with Movie Lens

```
def get_similar_movies( sim_df, movieid, topN = 5 ):
    """ get top-N similar movies given an input movie ...
        sim_df is the output of compute_movies_similarities
    """
    movies_df = pd.read_csv( "data/ml-100k/u.item", ... )
    movies_df = movies_df.iloc[:, :2]
    movies_df.columns = ['movieid', 'title']
    movies_df['similarity'] = sim_df.iloc[movieid - 1]
    movies_df.columns = ['movieid', 'title', 'similarity']
    top_n = # PUT YOUR CODE HERE
    return top_n
```

## Exercise 1: item-based recommendation with Movie Lens

```
def predict_rating (rating_df_pivoted, movie_sim_df, user_id):
    """ predict rating for a user and a movie
    """
    similar_movies = get_similar_movies (...)
    sim_ratings = []
    sim_scores = []
    for row2 in similar_movies[1:].itertuples():
        _, item_id_2, _, similarity = row2
        sim_rating = get_rating (...)
        if not numpy.isnan (sim_rating):
            sim_ratings.append (...)
            sim_scores.append (...)
    if len(sim_ratings) > 0:
        return numpy.dot(...) / numpy.sum (...)
    return numpy.nan
```



## Exercise 1: item-based recommendation with Movie Lens

```
def evaluate (rating_df_pivoted, movie_sim_df, num_ratings)
    """ predict ratings for the testing set and compute RMS
    """

    predicted_ratings = []
    true_ratings = []
    for row in ...[:num_ratings].itertuples():
        _, user_id, item_id, rating, _ = row
        predicted_rating = predict_rating (...)
        if not numpy.isnan (predicted_rating):
            predicted_ratings.append (predicted_rating)
            true_ratings.append (rating)

    rmse_val = rmse (...)
    return (rmse_val, len(predicted_ratings))
```

# Exercise 1: item-based recommendation with Movie Lens

Bonus points:

- play with the parameters to improve the RMSE
- compare to random
- ignore movies with less than  $k$  scores

## Exercise 2: MF with tensorflow

Given:

- data loader and matrix initialization (steps 1 and 2)
- default params settings, training and evaluation (steps 5 to 8)

Answer the following questions:

- Implement L1 and L2 losses (step 3)
- Implement L1 and L2 regularizations and evaluate (step 4)
- Bench the parameters around the default values

## Exercise 2: MF with tensorflow

Hints for step 3 (with L2 norm)

```
diff_op = # use tf.subtract and tf.add with  
          # result_values, mean_rating and  
          # rates as variables  
diff_op_squared = # use tf.square  
base_cost = # use tf.reduce_sum
```

## Exercise 2: MF with tensorflow

Hints for step 4 (with L2 norm)

```
l2_norm_sums = # use tf.add and tf.reduce_sum  
               # with U and P  
regularizer = # multiply l2_norm_sums and lamda  
# regularizer_cost = # add base_cost and regularizer
```

## Exercise 2: MF with tensorflow

Bonus points:

- How about another optimizer (e.g. AdamOptimizer)?
- Tune hyper parameters
- How should you set the number of epochs?